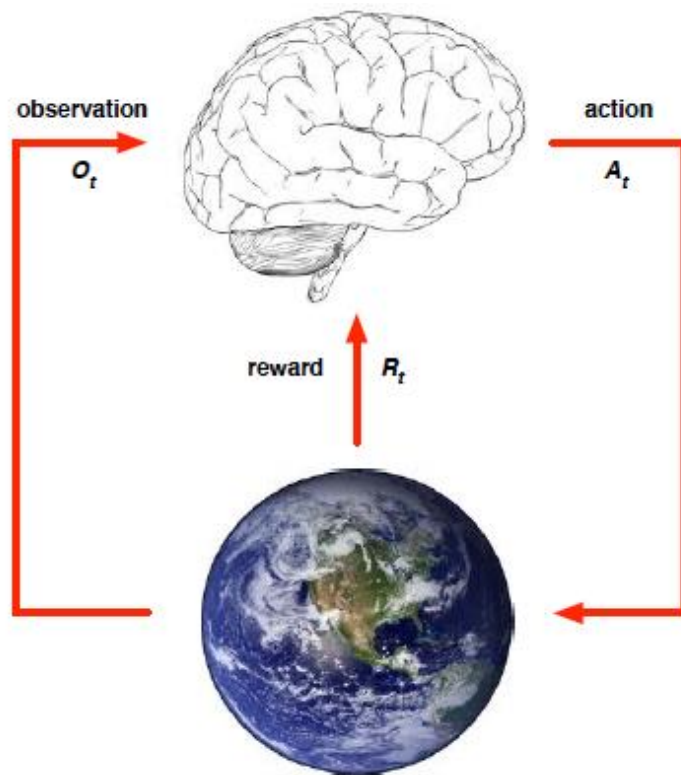


Lecture 2: MDP and Dynamic Programming

- Last lecture:
 - Course introduction
 - What's RL?
 - Broad applications of RL
 - Why RL?
- This lecture:
 - Basic components of RL: Reward, State, Policy, Model, Value function
 - The formal formulation of an RL problem as a Markov decision process
 - Making good decisions given a Markov decision process

The RL Problem

“learn to *make good sequences of decisions through trail-and-errors*”



- ❑ At each step t the agent:
 - ❑ Executes action A_t
 - ❑ Receives observation O_t
 - ❑ Receives scalar reward R_t
- ❑ The environment:
 - ❑ Receives action A_t
 - ❑ Emits observation O_{t+1}
 - ❑ Emits scalar reward R_{t+1}

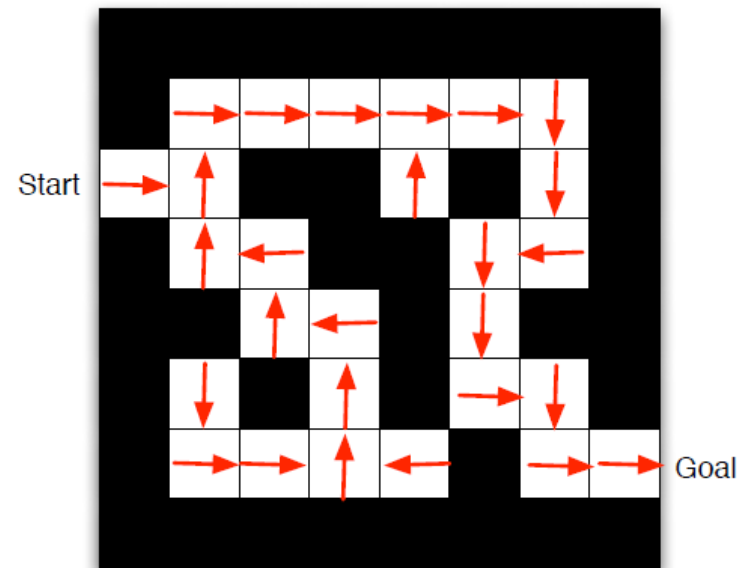
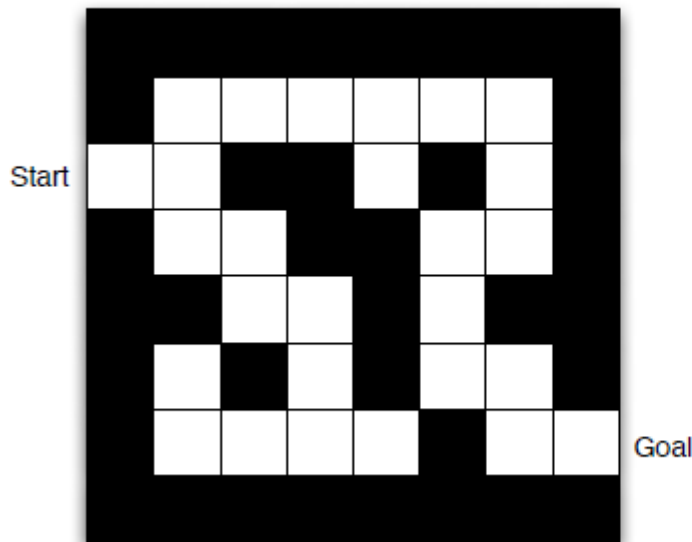
Goal: learn a policy (*i.e.*, a mapping from observations to actions) to maximise total future reward

Elements of RL Problems - Policy

□ **Policy**: an agent's behaviour function, i.e., a mapping from state to action

□ Deterministic policy: $a = \pi(s)$

□ Stochastic policy: $\pi(a|s) = \mathbb{P}[A_t = a | S_t = s]$




Elements of RL Problems - Policy

□ **Policy**: an agent's behaviour function, i.e., a mapping from state to action

□ Deterministic policy: $a = \pi(s)$

□ Stochastic policy: $\pi(a|s) = \mathbb{P}[A_t = a|S_t = s]$

s_1	s_2	s_3	s_4	s_5	s_6	s_7
						

- For the Mars rover example [7 discrete states (location of rover); 2 actions: Left or Right], how many deterministic policies are there?


$$2 / 14 / 7^2 / 2^7$$

Elements of RL Problems - Model

- ❑ **Model**: A model predicts what the environment will do next, i.e., agent's representation of the environment
- ❑ P predicts the next state
- ❑ R predicts the next (immediate) reward

$$\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$$

$$\mathcal{R}_s^a = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$$




Elements of RL Problems – Value Function


□ **Value functions:** how good is each state and/or action

- Value function is a prediction of future reward
- Used to evaluate the goodness/badness of states
- And therefore used to select between actions

$$v_{\pi}(s) = \mathbb{E}_{\pi} [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s]$$




Categorizing RL Algorithms



[from David Silver's course]



Categorizing RL Algorithms



<https://blog.csdn.net/gsw404>

Deep RL

Formulation of an RL Problem



- Markov Process
- Markov Reward Process (MRP)
- Markov Decision Process (MDP)
- Evaluation/Prediction and Improvement/Control in MDP

Recall: Markov Property

- A **Markov state** contains all useful information from the history, i.e., future is independent of past given present

A state S_t is **Markov** if and only if

$$\mathbb{P}[S_{t+1} \mid S_t] = \mathbb{P}[S_{t+1} \mid S_1, \dots, S_t]$$

- **Markov Process** or Markov Chain

- Sequence of random states with Markov property
- P is dynamics/transition model that specifies

$$p(s_{t+1} = s' \mid s_t = s)$$


- no rewards, no actions
- P can be expressed as a matrix

$$P = \begin{pmatrix} P(s_1|s_1) & P(s_2|s_1) & \cdots & P(s_N|s_1) \\ P(s_1|s_2) & P(s_2|s_2) & \cdots & P(s_N|s_2) \\ \vdots & \vdots & \ddots & \vdots \\ P(s_1|s_N) & P(s_2|s_N) & \cdots & P(s_N|s_N) \end{pmatrix}$$

Markov Process

s_1	s_2	s_3	s_4	s_5	s_6	s_7
						

The Mars rover problem



$$P = \begin{pmatrix} 0.6 & 0.4 & 0 & 0 & 0 & 0 & 0 \\ 0.4 & 0.2 & 0.4 & 0 & 0 & 0 & 0 \\ 0 & 0.4 & 0.2 & 0.4 & 0 & 0 & 0 \\ 0 & 0 & 0.4 & 0.2 & 0.4 & 0 & 0 \\ 0 & 0 & 0 & 0.4 & 0.2 & 0.4 & 0 \\ 0 & 0 & 0 & 0 & 0.4 & 0.2 & 0.4 \\ 0 & 0 & 0 & 0 & 0 & 0.4 & 0.6 \end{pmatrix}$$

e.g., Sample episodes starting from s_4

- $s_4, s_5, s_6, s_7, s_7, s_7, \dots$
- $s_4, s_4, s_5, s_4, s_5, s_6, \dots$
- $s_4, s_3, s_2, s_1, \dots$

Markov Reward Process (MRP)

□ **Markov Reward Process** is a Markov Process with rewards


□ P is dynamics/transition model that specifies

$$p(s_{t+1} = s' | s_t = s)$$

□ R is a reward function $R(s_t = s) = \mathbb{E}[r_t | s_t = s]$

□ Discount factor $\gamma \in [0, 1]$

□ No actions



Reward: +1 in s_1 , +10 in s_7 , 0 in all other states

Return & Value Function

- Definition of **Horizon** (H)

- Number of time steps in each episode
- Can be infinite or finite

- Definition of **Return**

- Discounted sum of rewards from time step t to horizon H

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \cdots + \gamma^{H-1} r_{t+H-1}$$

- Definition of **State Value Function** $V(s)$

- Expected return from starting in state s

$$V(s) = \mathbb{E}[G_t | s_t = s] = \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \cdots + \gamma^{H-1} r_{t+H-1} | s_t = s]$$

- ❑ Mathematically convenient
 - ❑ avoid infinite returns and values
- ❑ Model humans' behaviors
 - ❑ $\gamma = 0$: only care about immediate reward
 - ❑ $\gamma = 1$: future reward is with the same importance

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \cdots + \gamma^{H-1} r_{t+H-1}$$

Computing the Value of MRP


□ MRP value function satisfies the **Bellman Equation**

$$V(s) = \mathbb{E}[G_t | s_t = s] = \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \cdots + \gamma^{H-1} r_{t+H-1} | s_t = s]$$

$$V(s) = \underbrace{R(s)}_{\text{Immediate reward}} + \underbrace{\gamma \sum_{s' \in S} P(s'|s) V(s')}_{\text{Discounted sum of future rewards}}$$


Computing the Value of MRP

$$V(s) = \underbrace{R(s)}_{\text{Immediate reward}} + \underbrace{\gamma \sum_{s' \in S} P(s'|s) V(s')}_{\text{Discounted sum of future rewards}}$$



Computing the Value of MRP

$$V(s) = \underbrace{R(s)}_{\text{Immediate reward}} + \underbrace{\gamma \sum_{s' \in S} P(s'|s) V(s')}_{\text{Discounted sum of future rewards}}$$



Computing the Value of MRP

- For finite state MRP, we can express $V(s)$ in a matrix form

$$\begin{pmatrix} V(s_1) \\ \vdots \\ V(s_N) \end{pmatrix} = \begin{pmatrix} R(s_1) \\ \vdots \\ R(s_N) \end{pmatrix} + \gamma \begin{pmatrix} P(s_1|s_1) & \cdots & P(s_N|s_1) \\ P(s_1|s_2) & \cdots & P(s_N|s_2) \\ \vdots & \ddots & \vdots \\ P(s_1|s_N) & \cdots & P(s_N|s_N) \end{pmatrix} \begin{pmatrix} V(s_1) \\ \vdots \\ V(s_N) \end{pmatrix}$$

$$V = R + \gamma PV$$

$$V - \gamma PV = R$$

$$(I - \gamma P)V = R$$

$$V = (I - \gamma P)^{-1}R$$

- There are many iterative methods for large MRPs, e.g.
 - Dynamic programming
 - Monte-Carlo evaluation
 - Temporal-Difference learning

Markov Decision Process (MDP)

□ **Markov Decision Process** is Markov Reward Process with actions

□ P is dynamics/transition model for each action that specifies

$$P(s_{t+1} = s' | s_t = s, a_t = a)$$

□ R is a reward function $R(s_t = s, a_t = a) = \mathbb{E}[r_t | s_t = s, a_t = a]$

□ Discount factor $\gamma \in [0, 1]$

□ MDP is a tuple: (S, A, P, R, γ)

$$P(s'|s, a_1) = \begin{pmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{pmatrix} \quad P(s'|s, a_2) = \begin{pmatrix} 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$$

The transition matrix for the Mars rover problem (a_1 means moving left, and a_2 means moving right)

- Policy specifies what action to take in each state
 - Can be deterministic or stochastic
 - Usually is a distribution over actions given states

$$\pi(a|s) = P(a_t = a | s_t = s)$$

- Given an MDP and a policy, then

$$R^\pi(s) = \sum_{a \in A} \pi(a|s) R(s, a)$$


$$P^\pi(s'|s) = \sum_{a \in A} \pi(a|s) P(s'|s, a)$$

- **State-Action Value Q** for a policy

$$Q^\pi(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V^\pi(s')$$

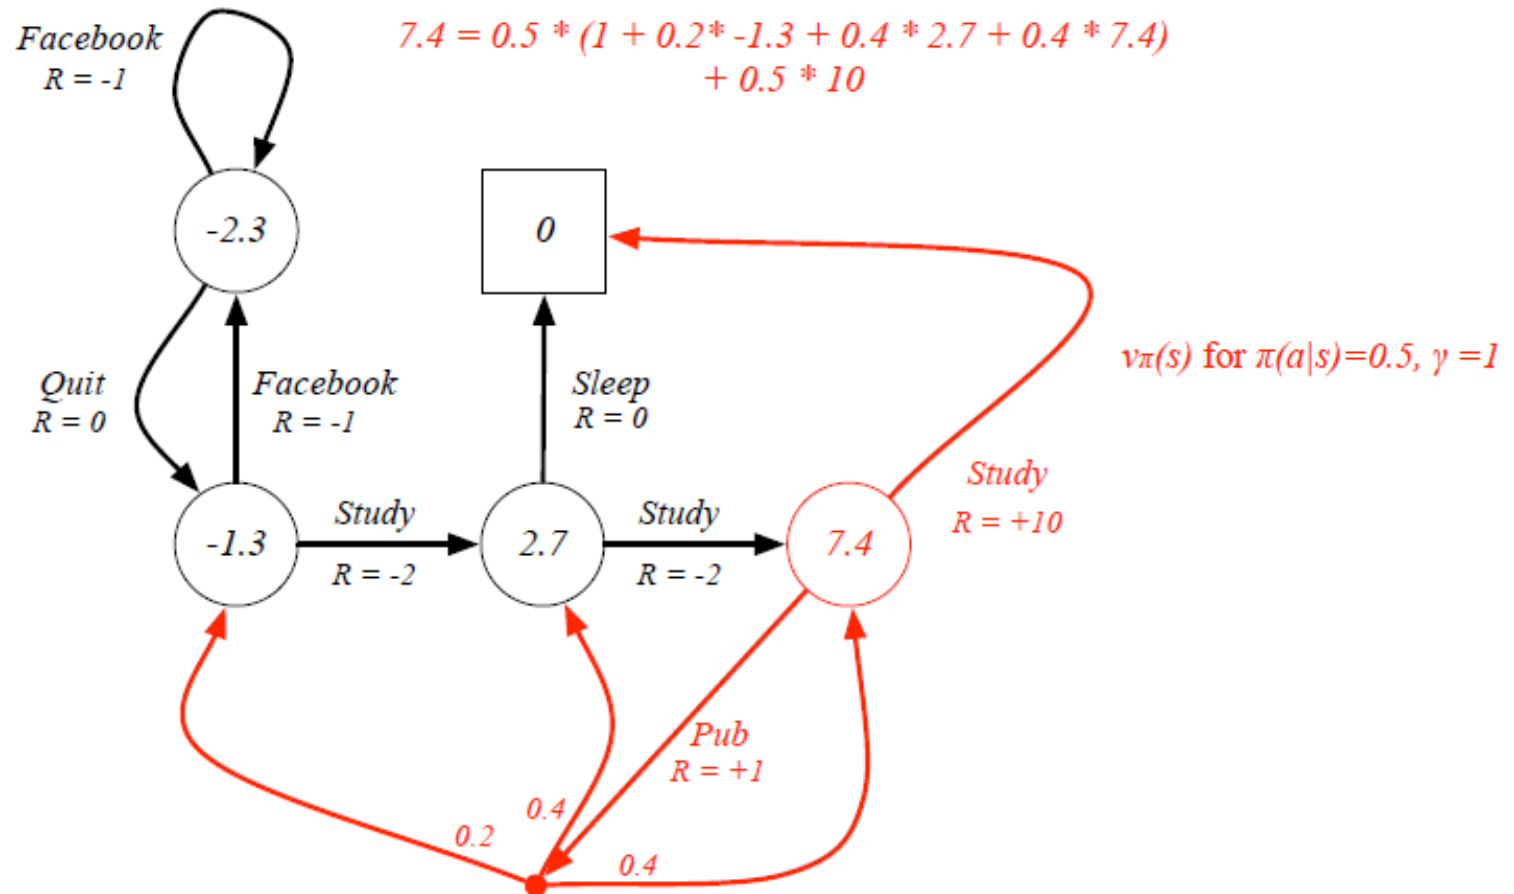
Computing the Value of MDP

- The value of a state is the value of expected next state plus the reward expected along the way



Computing the Value of MDP

- The value of a state is the value of expected next state plus the reward expected along the way



- ❑ Compute the **optimal policy**

$$\pi^*(s) = \arg \max_{\pi} V^{\pi}(s)$$

- ❑ There exists a unique optimal value function, which specifies the best possible performance in the MDP
- ❑ Optimal policy for an MDP in an infinite horizon problem is deterministic, but not necessarily unique
- ❑ One option is searching to compute best policy
- ❑ Number of deterministic policies is $|A|^{|S|}$
- ❑ Policy iteration is generally more efficient than enumeration

What's Dynamic Programming (DP)?

- ❑ *Dynamic*: sequential or temporal component to the problem
- ❑ *Programming*: optimizing a “program”, i.e. a policy
- ❑ A method for solving complex problems by
 - ❑ breaking them down into subproblems
 - ❑ combining solutions of subproblems
- ❑ DP is a general solution method for problems with two properties:
 - ❑ Optimal solution can be decomposed into subproblems
 - ❑ Subproblems recur many times and solutions can be cached and reused
- ❑ MDP satisfy both properties
 - ❑ Bellman equation gives recursive decomposition
 - ❑ Value function stores and reuses solutions

$$V(s) = \underbrace{R(s)}_{\text{Immediate reward}} + \underbrace{\gamma \sum_{s' \in S} P(s'|s) V(s')}_{\text{Discounted sum of future rewards}}$$

Solving MDP using DP



- ❑ DP assumes full knowledge of the MDP for planning
- ❑ DP is an iterative solution method to MDP
 - ❑ Policy Iteration (PI)
 - ❑ Value Iteration (VI)
- ❑ PI iterates between the following processes
 - ❑ Policy **evaluation** (prediction): Estimate/predict the expected rewards from following a given policy
 - ❑ Policy **improvement** (control): find a better policy
- ❑ VI iterates between the estimation of value functions and policy optimization, without explicit policy

Value Function for MDP



- The state-value function $v(s)$ of an MDP is the expected return starting from state s , and following policy π

$$v^\pi(s) = \mathbb{E}_\pi[G_t | s_t = s]$$

where $G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$

- The action-value function $q(s,a)$ is the expected return starting from state s , taking action a , and then following policy π

$$q^\pi(s, a) = \mathbb{E}_\pi[G_t | s_t = s, A_t = a]$$

- We have the relation

$$v^\pi(s) = \sum_{a \in A} \pi(a|s) q^\pi(s, a)$$

Bellman Expectation Equation



- The state-value function can be decomposed into immediate reward plus discounted value of the successor state,

$$v^{\pi}(s) = E_{\pi}[R_{t+1} + \gamma v^{\pi}(s_{t+1}) | s_t = s]$$

- The action-value function can similarly be decomposed

$$q^{\pi}(s, a) = E_{\pi}[R_{t+1} + \gamma q^{\pi}(s_{t+1}, A_{t+1}) | s_t = s, A_t = a]$$

Bellman Expectation Equation for V and Q



$$v^{\pi}(s) = \sum_{a \in A} \pi(a|s) q^{\pi}(s, a)$$


$$q^{\pi}(s, a) = R_s^a + \gamma \sum_{s' \in S} P(s'|s, a) v^{\pi}(s')$$

Thus


$$v^{\pi}(s) = \sum_{a \in A} \pi(a|s) (R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) v^{\pi}(s'))$$

$$q^{\pi}(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) \sum_{a' \in A} \pi(a'|s') q^{\pi}(s', a')$$

Backup Diagram for V and Q





$$v^\pi(s) = \sum_{a \in A} \pi(a|s) (R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) v^\pi(s'))$$



$$q^\pi(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) \sum_{a' \in A} \pi(a'|s') q^\pi(s', a')$$

Policy Iteration (PI)

- Set $i = 0$
- Initialize $\pi_0(s)$ randomly for all states s
- While $i == 0$ or $\|\pi_i - \pi_{i-1}\|_1 > 0$ (L1-norm, measures if the policy changed for any state):
 - $V^{\pi_i} \leftarrow$ MDP V function policy **evaluation** of π_i
 - $\pi_{i+1} \leftarrow$ Policy **improvement**
 - $i = i + 1$



Policy Evaluation

- ❑ Objective: evaluate a given policy π for a MDP
- ❑ Output: the value function under policy π
- ❑ Solution: iteration on Bellman expectation backup
- ❑ Algorithm: Synchronous backup


- 1 At each iteration $t+1$
update $v_{t+1}(s)$ from $v_t(s')$ for all states $s \in \mathcal{S}$ where s' is a successor state of s

$$v_{t+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) (R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) v_t(s'))$$

- ❑ Convergence: $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v^\pi$


□ Example 4.1 in the Sutton RL textbook.

- Actions leading out of the grid leave state unchanged
- Reward is -1 until the terminal state is reached
- Agent follows uniform random policy



$$v_{t+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) (R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) v_t(s'))$$

Policy Evaluation




Policy Evaluation




$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0




$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0



$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0



optimal policy

Policy Improvement

□ Consider a deterministic policy $a = \pi(s)$

□ We improve the policy through

$$\pi'(s) = \arg \max_a q^\pi(s, a)$$

□ This improves the value from any state s over one step

$$q^\pi(s, \pi'(s)) = \max_{a \in \mathcal{A}} q^\pi(s, a) \geq q^\pi(s, \pi(s)) = v^\pi(s)$$

□ It therefore improves the value function $v_{\pi'}(s) \geq v_\pi(s)$

$$\begin{aligned} v_\pi(s) &\leq q_\pi(s, \pi'(s)) = \mathbb{E}_{\pi'} [R_{t+1} + \gamma v_\pi(S_{t+1}) \mid S_t = s] \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma q_\pi(S_{t+1}, \pi'(S_{t+1})) \mid S_t = s] \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma R_{t+2} + \gamma^2 q_\pi(S_{t+2}, \pi'(S_{t+2})) \mid S_t = s] \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma R_{t+2} + \dots \mid S_t = s] = v_{\pi'}(s) \end{aligned}$$

PI as Bellman Operations

□ **Bellman backup operator** B^π for a particular policy is defined as

$$B^\pi V(s) = R^\pi(s) + \gamma \sum_{s' \in S} P^\pi(s'|s) V(s)$$

$$V^\pi(s) = R^\pi(s) + \gamma \sum_{s' \in S} P^\pi(s'|s) V^\pi(s')$$

□ Policy evaluation amounts to computing the **fixed point** of B^π

□ To do policy evaluation, repeatedly apply operator until V stops changing

$$V^\pi = B^\pi B^\pi \dots B^\pi V$$

□ To do policy improvement,

$$\pi_{k+1}(s) = \arg \max_a R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V^{\pi_k}(s')$$

Bellman Optimality Equation



- The optimal value function v^* are reached by the **Bellman optimality equations**

$$v^*(s) = \max_a q^*(s, a)$$

$$q^*(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) v^*(s')$$

thus

$$v^*(s) = \max_a R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) v^*(s')$$

$$q^*(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) \max_{a'} q^*(s', a')$$

- ❑ Policy iteration computes optimal value and policy
- ❑ Value iteration is another technique
 - ❑ Maintain optimal value of starting in a state s if having a finite number of steps k left in the episode
 - ❑ Iterate to consider longer and longer episodes
- ❑ In other word, we assume we know the solution to subproblems and then find the optimal solution by one-step lookahead

$$v(s) \leftarrow \max_{a \in \mathcal{A}} \left(R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) v(s') \right)$$

Value Iteration (VI)

Example: Shortest Path

g			

Problem

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

V_1

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1

V_2

0	-1	-2	-2
-1	-2	-2	-2
-2	-2	-2	-2
-2	-2	-2	-2

V_3

0	-1	-2	-3
-1	-2	-3	-3
-2	-3	-3	-3
-3	-3	-3	-3

V_4

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-4
-3	-4	-4	-4

V_5

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-5

V_6

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-6

V_7

Value Iteration (VI)

- ① Objective: find the optimal policy π
- ② Solution: iteration on the Bellman optimality backup
- ③ Value Iteration algorithm:

- ① initialize $k = 1$ and $v_0(s) = 0$ for all states s
- ② For $k = 1 : H$
 - ① for each state s

$$q_{k+1}(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) v_k(s')$$

$$v_{k+1}(s) = \max_a q_{k+1}(s, a)$$

- ② $k \leftarrow k + 1$
- ③ To retrieve the optimal policy after the value iteration:

$$\pi(s) = \arg \max_a R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) v_{k+1}(s')$$

Value vs Policy Iteration



- ❑ Policy iteration includes: **policy evaluation** + **policy improvement**, and the two are repeated iteratively until policy converges
- ❑ Value iteration includes: **finding optimal value function** + **one policy extraction**. There is no repeat of the two because once the value function is optimal, then the policy out of it should also be optimal (i.e. converged).
- ❑ Finding optimal value function can also be seen as a combination of policy improvement (due to max) and truncated policy valuation (the reassignment of $v(s)$ after just one sweep of all states regardless of convergence).

Will DP Converge?



□ Bellman backup operator is a contraction if discount factor when $\gamma < 1$

□ Contraction Operator

- Let O be an operator, and $|x|$ denote (any) norm of x
- If $|OV - OV'| \leq |V - V'|$, then O is a contraction operator

□ If apply it to two different value functions, distance between value functions shrinks after applying Bellman equation to each

- Let $\|V - V'\| = \max_s |V(s) - V'(s)|$ be the infinity norm

$$\begin{aligned}
 \|BV_k - BV_j\| &= \left\| \max_a \left(R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_k(s') \right) - \max_{a'} \left(R(s, a') + \gamma \sum_{s' \in S} P(s'|s, a') V_j(s') \right) \right\| \\
 &\leq \max_a \left\| \left(R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_k(s') - R(s, a) - \gamma \sum_{s' \in S} P(s'|s, a) V_j(s') \right) \right\| \\
 &= \max_a \left\| \gamma \sum_{s' \in S} P(s'|s, a) (V_k(s') - V_j(s')) \right\| \\
 &\leq \max_a \left\| \gamma \sum_{s' \in S} P(s'|s, a) \|V_k - V_j\| \right\| \\
 &= \max_a \left\| \gamma \|V_k - V_j\| \sum_{s' \in S} P(s'|s, a) \right\| \\
 &= \gamma \|V_k - V_j\|
 \end{aligned}$$

Summary of DP methods

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Iterative Policy Evaluation
Control	Bellman Expectation Equation + Greedy Policy Improvement	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

- Algorithms are based on state-value function $v_{\pi}(s)$ or $v_*(s)$
- Complexity $O(mn^2)$ per iteration, for m actions and n states
- Could also apply to action-value function $q_{\pi}(s, a)$ or $q_*(s, a)$
- Complexity $O(m^2n^2)$ per iteration

Asynchronous DP

- ❑ DP methods described so far used *synchronous* backups
 - ❑ i.e. all states are backed up in parallel
- ❑ *Asynchronous* DP backs up states individually, in any order
 - ❑ For each selected state, apply the appropriate backup
 - ❑ Can significantly reduce computation
 - ❑ Guaranteed to converge if all states continue to be selected
- ❑ Three simple ideas for asynchronous DP:
 - ❑ **In-place DP**: only stores one copy of value function
 - ❑ **Prioritised sweeping**: backup the state with the largest remaining Bellman error
$$\left| \max_{a \in \mathcal{A}} \left(\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v(s') \right) - v(s) \right|$$
 - ❑ **Real-time DP**: Use agent's experience to guide the selection of states

Full-Width Backups

- ❑ DP uses **full-width backups**
 - ❑ For each backup (sync or async), every successor state and action is considered
 - ❑ DP is effective for medium-sized problems (millions of states), but suffers Bellman's curse of dimensionality for large DP, as number of states grows exponentially with number of state variables
- ❑ Sample backups use sample rewards and sample transitions, instead of reward function and transition dynamics
 - ❑ **Model-free**: no advance knowledge of MDP required
 - ❑ Breaks the curse of dimensionality through sampling
 - ❑ Cost of backup is constant, independent of $n = |\mathcal{S}|$

- ❑ Define MP, MRP, MDP, Bellman operator, contraction, model
- ❑ Know how to compute value function, optimal policy
- ❑ Be able to implement Value Iteration and Policy Iteration
- ❑ Give pros and cons of different policy evaluation approaches
- ❑ Be able to prove contraction properties